Draft Technical Notes on BayesDB

August 13, 2015

1 Overview: Generative Population Models and the BayesDB Minimal Modeling Language

BQL programs are executed against a set of ambient data *generative population models* (GPMs), with one GPM per table in the database. These generative population models are provided by BayesDB based on modeling schemas and inference programs written in the BayesDB Minimal Modeling Language (MML). Each GPM induces a random table $\mathbf{X}$ with a finite number of columns and an infinite number of rows. Some of the rows are observed and represent the actual dataset stored in BayesDB. Each cell $\mathbf{X_{ij}}$ in the table represents a random variable, and may have one or more realizations $x_{ij}^{(t)}$.

A GPM $\mathcal{G}$ can simulate and evaluate the probability of an arbitrary collection of cells in the random table, for either observed or hypothetical rows.

> Note sure how to formalize this in a tabular form.

In its simplest form a GPM $\mathcal{G}$ is a fixed object that fully characterizes a data generating process, and observing generated members has no bearing on the actual structure.

Example 1.1 *A Pretty Normal Generative Population Model*

Consider a GPM $\mathcal{G}$ which simulates D-dimensional members of a statistical population according to a mixture of two normal distributions such that $p(x) = w_1\mathcal{N}(x; \mu_1, \Sigma_1) + w_2\mathcal{N}(x; \mu_2, \Sigma_2)$. The structure of the GPM can be thought of as:

- *The generative population model class $\mathcal{M}$ from which this GPM $\mathcal{G}$ is derived. Here $\mathcal{M} = \{M_i^\theta\text{: mixture of i normals with parameters }\theta\}$.*

- *The member of the class corresponding to $\mathcal{G}$. In this case, $\mathcal{G} = M_2^{\theta^*}$ is indexed by $i^* = 2$ (the number of components) and parameters $\theta^* = \{w_1, w_2, \mu_1, \mu_2, \Sigma_1, \Sigma_2\}$.*

A population with N members $\{X^{(i)} := (X_1^{(i)}, \ldots, X_D^{(i)})\}_{i=1}^N$ generated by $\mathcal{G}$ is represented in BayesDB as a data table $\mathbf{X}$ with D columns and N rows.

The GPM in Example 1.1 is not particularly interesting. The detailed statistical characteristics of the GPM are fixed, but such information is rarely known a-priori [1].

In typical applications we are interesting in *learning the structure of the GPM from data*. In fully generality, this might be learning the specific parameters θ given a fixed generative process M_i, or learning the generative process M_i from some fixed model class $\mathcal{M}$, or even learning the model class itself $\mathcal{M}$ from a set of competing model classes $\{\mathcal{M}_k\}$. The Minimal Modeling Language provides a framework for jointly learning the structure of a collection of GPM $\{\mathcal{G}_i\}$ that interact with another to produce a statistical population.

A GPM $\mathcal{G}$ which can be learned from data is referred to an *adaptive generative population model*. An adaptive GPM $\mathcal{G}$ is probabilistic model defined over a space of GPMs. We will see in Section 3 that a adaptive GPM is a GPM in its own right and can answer the same questions about the statistical population **X**. The MML includes a *default adaptive generative population model* which defines a distribution over a collection of *GPMs from a default class* using a hierarchal, semi-parametric Bayesian model derived from CrossCat.

Example 1.2 *What Exactly Is CrossCat?*

A cross-categorization of a data table **X** *with D columns and N is a partition of the columns* $(X_1, \ldots, X_D)$ *into blocks called views. In each view, we have a partition of the rows* $(X_1^i, \ldots, X_D^i)_{i=1}^N$ *into blocks called categories.*


Figure 1: Some possible cross-categorizations of a table with 3 rows and 4 columns. The CrossCat model is very complex and contains several parameters such as the row partition, the column partition, and component parameters (which are shared by cells of the same color). Any two cells with a different color are marginally independent conditioned on the entire latent structure.

The collection of all possible column/row partitions j and associated component parameters θ of **X** *defines a GPM class:*

$$\mathcal{M}_X = \{CC_j^\theta : \text{cross-cat of } \mathbf{X} \text{ with partition } j \text{ and component model params } \theta\}.$$

A member $\mathcal{G}$ of the CrossCat GPM class is a particular cross- categorization and component parameters $CC_j^{\theta^*}$ from the model class. Since we do not know j

---

[1]One plausible scenario in which a 'fixed' GPM of this form can arise is if a user learns about the structure externally from the MML. Once the GPM is ported to BayesDB, the user decides no further inference or learning should occur.

or θ^*, we instead place distribution over all GPMS with this common structure. The object which carries this distribution over $\mathcal{M}_X$ is a adaptive generative population model, whose population members are GPMs that generate the $\mathbf{X}$!

Unlike Example 1.1, the generation process is inherently tied to the statistical population it has produced $\mathbf{X}$. Observing more generated members $\mathbf{X} \rightarrow \mathbf{X}'$ means the GPM class under consideration grows from $\mathcal{M}_X$ to $\mathcal{M}_{X'}$ (both the number of possible cross-categorizations j and the dimensionality of the vector of latents θ increase).

> The reason I am being so pedantic about this construction is to explicate various design choices when simulating, sampling, and running posterior inference. In particular we are going to have to break down the latent structure into 'global' vs 'row-related' latents to differentiate between observed and hypothetical members.

The default adaptive GPM is not obliged to model all columns jointly by positing some GPM from the CrossCat class. The MML includes modeling instructions for describing dependence constraints, and specifying a compositional, directed-acyclic network of arbitrary GPMs implemented by a user. It also includes inference instructions for initializing adaptive GPMs and performing updates of the posterior distribution on GPMs.

The MML is thus a complete (but sub- Turing) probabilistic programming language that interoperates with queries written in BQL. The set of primitive GPMs can be extended by loading foreign GPMs written in VentureScript, and can be used to transparently integrate models from disparate probabilistic programming languages such as Stan and statistical computing environments such as R. Because BQL programs only depend on the statistical data types of the columns, the underlying modeling technologies, GPM assumptions, and inference tactics can be changed without invalidating end-user data analysis workflows.

2 Interface for Generative Population Models

A *generative population model* is a probabilistic model of a data generating process and a statistical population that it has produced. It is informative to think of this statistical population as an infinite table $\mathbf{X}$ with a finite number of columns D (representing attributes) and an infinite number of rows (representing members of the population). A finite number of rows $r_i \in \{1, \ldots, N\}$ are actually observed, so $\mathbf{X}$ is stored in BayesDB as an $N \times D$ table. Hypothetical members are indexed by drawing a row index $r^* \sim U[0, 1]$, which guarantees their uniqueness.

Discussion 2.1 *Technical Aside On IID Rows*

The classical statistical framework generally assumes that all members of a population are independent and identically distributed. While a particular generator is free to make any assumptions about its population, the IID assumption is a strong one. The default generator (and almost any generator being learned under the Bayesian framework) will usually assume that the rows are not IID but exchangeably coupled. Di Finetti's theorem guarantees that there exists a mixture of random measures $\{\mathcal{G}_\alpha\}$ (which in most cases is indexed by the latent state of the GPM) where the population is conditionally IID.

$$p(\mathbf{X}) = \int_\alpha (\Pi_{i=1}^N p(\mathbf{X}_i | \mathcal{G}_\alpha)) d\mathcal{Q}(\mathcal{G}_\alpha)$$

The Di Finetti measure $\mathcal{Q}$ is the object which defines a distribution over the (fixed) measures (GPMs) that produce a statistical population. In this sense, $\mathcal{Q}$ represents the adaptive GPM.

Generative population models provide the primitive statistical inferences that are used by BayesDB to implement inferential queries in BQL. To support all of BQL, a GPM must provide the ability to **sample from** and **evaluate the log density** of all possible marginal distributions subject to equality constraints for arbitrary subsets of variables. The interface for *initializing* and *querying* GPMs is outlined below.

1. $\mathcal{G}$ = `initialize(` `schema` = Λ`)`

   Initialize a GPM with the given schema, and return the resulting GPM $\mathcal{G}$.

   Each GPM is initialized with a *schema* $\Lambda =$(`typed-outputs`,`typed-inputs`, `body`). The `typed-outputs` component species the column indexes and statistical types of each column that the GPM will be responsible for generating. The `typed-inputs` component specifies the indexes and statistical types of columns that the data generator can read from. The `body` is an opaque binary that contains any generator specific configuration information, such as a probabilistic program.

2. $\{\vec{s}_i\}_{i=1}^N$ = `simulate(`$\mathcal{G}$, `targets` = $\{(r_i^t, c_i^t)\}_{i=1}^{|T|}$, `givens` = $\{(r_i^g, c_i^g, x_{(r_i^g, c_i^g)})\}_{i=1}^{|G|}$, `samples` = N`)`

   Generate N samples from the specified conditional distribution.

   $$\{\vec{s}_i\}_{i=1}^N \sim p(\{\mathbf{X}_{(r_i^t, c_i^t)}\} | \{\mathbf{X}_{(r_i^g, c_i^g)} = x_{(r_i^g, c_i^g)}\}, \mathcal{G})$$

   Note that each $\vec{s}_i$ is a $|T|$ dimensional vector

3. $\log p$ = `logpdf(`$\mathcal{G}$, `targets` = $\{(r_i^t, c_i^t, x_{(r_i^t, c_i^t)})\}_{i=1}^{|T|}$, `givens` = $\{(r_i^g, c_i^g, x_{(r_i^g, c_i^g)})\}_{i=1}^{|G|}$

   Evaluate the log probability density of the specified conditional distribution at a target point.

$$\log p = p(\{\mathbf{X}_{(r_i^t, c_i^t)} = x_{(r_i^t, c_i^t)}\} | \{\mathbf{X}_{(r_i^g, c_i^g)} = x_{(r_i^g, c_i^g)}\}, \mathcal{G})$$

For a BayesDB table $\mathbf{X}$ with D columns and N rows, the allowed values for the indexes are

$$c_i \in \{1, \ldots, D\}$$

$$r_i \in \{1, \ldots, N\} \cup [0, 1]$$

The implications of this restriction are that the MML does not have a notion of a hypothetical column. The row indexes in $\{1, \ldots, N\}$ correspond to observed rows, while $r_i \in [0, 1]$ are hypothetical members whose uniqueness is guaranteed by simulating from the uniform distribution on the unit interval.

While the requirement that the GPM should respond to questions about arbitrary cells in the random table may appear overwhelmingly complex, the structure of most GPMs induces several internal simplifications.

3   Extended Interface for Adaptive Generative Population Models

As outlined in Section 1, some GPMs can be learned from an observed statistical population $\mathbf{X}$. Learning the GPM structure will be based on approximate probabilistic inference in a *adaptive generative population model* $\mathcal{G}$ which is a probabilistic model defined over a space of GPMs $\{\mathcal{G}_\alpha\}$. The GPMs $\mathcal{G}_\alpha$ typically share structure (such as belonging to the same parametric class) but this is not a formal restriction.

A key insight about a adaptive GPM $\mathcal{G}$ is that it can *answer the same questions about the underlying statistical population* $\mathbf{X}$ that any one of its GPMs $\mathcal{G}_\alpha$ can. These questions are answered achieved by an appropriate sequence of sampling, conditioning, and marginalization operations over the space of GPMs $\mathcal{G}_\alpha$. Moreover, a adaptive GPM $\mathcal{Q}$ can itself define a distribution over a space of adaptive GPMs, up to an infinitely deep level of recursion.

Thus far, all adaptive GPMs in BayesDB have been based on *Markov chain* methods. These adaptive GPMs internally maintain a sample from an approximate posterior over GPMS, and provide a Markov chain transition operator that updates the sample stochastically. Some Markov chains are *asymptotically Bayesian* in the sense that the distribution that results from sequences of T updates converges to the posterior over GPMs as T tends to infinity.

The extended interface for adaptive GPMs based on Markov chains is outlined below:

1. $\mathcal{G}$ = initialize(schema = Λ)

Initializes a adaptive GPM with the given schema Λ = (`typed-outputs`, `typed-inputs`, `body`), and returns an arbitrary initialization $\theta_{\mathcal{G}}^0$ and empty tabular data store $\mathbf{X}_{\mathcal{G}}$.

> If there are columns in `typed-outputs` or `typed-inputs` for which no generative procedure is defined in `body`, then the default GPM will be used to jointly model those columns. This is required because the MML needs a full generative model to query arbitrary patterns of population members. Otherwise several important queries would be ill-defined (for instance, simulating an output without conditioning on an input).

2. `incorporate(value = `$(r_j, c_j, x_{(r_j, c_j)})$`)`

   Records an observation $x_{(r_j, c_j)}$ in the cell $\mathbf{X}_{\mathbf{r_j}, \mathbf{c_j}}$. If there is no row r_j, a new one will be created.

   Errors result from overwriting existing cells, or inserting values $x_{(r_j, c_j)}$ incompatible with the meta-schema Λ (for example providing a wrong data type).

3. `remove(`(r_j, c_j)

   Removes the observation $x_{(r_i, c_j)}$ stored in cell $\mathbf{X}_{ij}$.

4. `infer(program = `$\mathcal{P}$`)`

   Simulate an internal Markov chain transition operator $\mathcal{T}$ based on the inference procedure specified in the program $\mathcal{P}$.

   This inference program is typically designed in cohesion with the `body` specified in the schema Λ. Each run of `infer` improves the quality of the posterior sample.

Some Markov chain GPMs are *asymptotically Bayesian*:

$$\lim_{t \to \infty} D_{KL}(p(\theta_{\mathcal{G}} | \mathbf{X}_{\mathcal{G}}) || p(\mathcal{T}^t(\theta_{\mathcal{G}}))) = 0$$

A sufficiently expressive asymptotically Bayesian Markov chain GPM may also be *asymptotically consistent* in the usual sense.

4   An Ensemble of Adaptive Generative Population Models

If BayesDB inferences were based on a single instance of a adaptive GPM, the inferences would arbitrarily suppress modeling uncertainty. The MML provides

an alternative constructor for adaptive GPMs that internally manages and ensemble of GPMs. In the language of Markov Chains, the ensemble represents a set of independent, parallel chains.

1. $\mathcal{Q}$ = initialize(schema = Λ, count = S)

   Initializes a collection of N GPMs with the schema Λ.

   $$\mathcal{Q} = (\{\mathcal{G}_s = (\theta_{\mathcal{G}}^{(s,0)})\}_{s=1}^{S}, \mathbf{X}_{\mathcal{Q}})$$

   Note that the schema and data store are shared across all GPMs in the ensemble.

The default ensemble GPM maintains a uniform distribution over all N GPMs and approximately marginalizes over the posterior on GPMs via the appropriate form of simple Monte Carlo.

1. simulate delegates to a randomly sampled Markov chain GPM $\mathcal{G}_k$ for each of the N samples in the returned set.

   $$\{\vec{s}_i\} = \cup_k \texttt{simulate}(\mathcal{G}_k, \ldots, \texttt{size} = 1) \text{ where } \mathcal{G}_k \sim U[\{\mathcal{G}_s\}]$$

2. logpdf forms a Monte Carlo estimate of the expected density using all of the GPMs in the ensemble.

   $$\log p = \frac{1}{S} \sum_s \texttt{logpdf}(\mathcal{G}_s)$$

Other functions in the GPM interface are defined analogously. More elaborate methods such as likelihood weighting of posterior samples from each chain are a possible extension.

Excercise 4.1 *Testing the Default Generative Population Model*

Consider a adaptive GPM $\mathcal{G}_N$ which defines a distribution over the space of mixture of Gaussian GPMs of the form seen in Example 1.1. Rather than fix the parameters as we did earlier, now place a prior over the number of Gaussian components i and the weights, means, and covariances θ. Also consider the default GPM $\mathcal{G}_{CC}$ from Example 1.2, which has a vague prior over the cross-categorizations and mixture parameters.

Suppose a population $\mathbf{X}$ is observed, and we are interested in comparing the two GPMs. If you know about the internals of CrossCat, it should be clear that the collection of GPMs that $\mathcal{G}_N$ can generate is a subset of the GPMs generated by $\mathcal{G}_{CC}$ (how?).

How can you invoke the GPM interface to test the performance of the default metamodel? (Hint: first define a set of "performance" metrics and a synthetic dataset). Does it make sense for $\mathcal{G}_N$ to outperform $\mathcal{G}_{CC}$ in certain regimes?

Can you think of other GPMs that are special members of the default GPM class (hint: Naive Bayes, . . .), and associated tests?

5   Characterizing Probabilistic Dependence

Many analysis subproblems correspond directly to invocations of the GPM interface. For example, missing or unknown values can be inferring an set of results via `simulate` and reducing these to a point estimate. MAP prediction can be approximated the single result with the highest likelihood under `logpdf`.

The GPM interface is also sufficient to implement generic mechanisms for characterizing context-specific conditional dependence and independence. This mechanism can be used to implement additional measures that are useful for data analysis, such as a context-sensitive measure of similarity.

Below are a collection of queries that can be answered using semantics from the GPM interface. Consider a GPM $\mathcal{G}$ ascribing to either the primitive or adaptive GPM interface.

1. $\{I_k\} = \texttt{CMI}(\mathcal{G}, A = \{(r_i^a, c_i^a)\}_{i=1}^{|A|}, B = \{(r_i^b, c_i^b)\}_{i=1}^{|B|}, C = \{(r_i^c, c_i^c)\}_{i=1}^{|D|}, D = \{(r_i^a, c_i^a, x_{(r_i^a, c_i^a)})\}_{i=1}^{|D|},\ \texttt{accuracy} = N,\ \texttt{size} = K)$

   CMI is an abbreviation of `conditional-mutual-information` $(A : B | C, D = d)$ under the given GPM. Let A, B, C, and D be subsets of the random variables in the infinite random table $\mathbf{X}$ induced by the GPM. An important special case in which CMI can be used to compute bivariate, unconditional mutual information for distinct variables will have $C, D = \emptyset$, $A = \{(c_1^a, r^*)\}$ and $B = \{(c_1^b, r^*)\}$, with $r^* \sim U[0, 1]$ and $c_1^a \neq c_1^b$.

   One implementation of `CMI`, applicable to all generators, is Monte Carlo estimation as follows.

   Algorithm 5.1 *First generate a collection of N samples from $(A, B, C) | D = d$, each corresponding to a unique member of the population.*

   $$r_i \sim U[0, 1]\ \textit{where}\ i \in \{0, \ldots, N - 1\}$$
   $$\{(\hat{a}_i, \hat{b}_i, \hat{c}_i)\} \sim (A, B, C) | D = d$$

   These samples can be used to form a Monte Carlo estimate of the desired conditional mutual information.

$$I(A:B|C, D=d) := \int \log \frac{p_{C|D=d}(c)p_{A,B,C|D=d}(a,b,c)}{p_{A,C|D=d}(a,c)p_{B,C|D=d}(b,c)} dF_{A,B,C|D=d}(a,b,c)$$

$$\approx \sum_{(a,b,c)\in\{(\hat{a}_i,\hat{b}_i,\hat{c}_i)\}} \frac{p_{C|D=d}(\hat{c}_i)p_{A,B,C|D=d}(\hat{a}_i,\hat{b}_i,\hat{c}_i)}{p_{A,C|D=d}(\hat{a}_i,\hat{c}_i)p_{B,C|D=d}(\hat{b}_i,\hat{c}_i)}$$

The output of `CMI` is set of K estimates $\{I_k\}$, where each estimate is created using N samples. This allows us to characterize uncertainty over the true value by computing Monte Carlo standard errors and other quantities of interest.

Note that the samples are generated by invoking

$$\texttt{simulate}(\mathcal{G}, \texttt{targets} = \{(r^{\cdot}c_i^a)\} \cup \{r^{\cdot}c_i^b\} \cup \{r^{\cdot}c_i^c\}, \texttt{givens} = \{(r_i^b, c_i^b, x_{(r_i^b,c_i^b)})\}, \texttt{size} = 1)$$

and the densities in the estimated are evaluated by invoking the corresponding `logpdf` command. An alternative implementation of Algorithm 5.1 would be to sample a single row index $r^* \sim U[0,1]$ and using N samples.